

centra2

March 20, 2024

```
[63]: #importing the necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
```

```
[64]: # loading my data
data=pd.read_csv('C:\\Users\\lynda\\Desktop\\apple_quality.csv')
data.head()
```

```
[64]:
```

	A_id	Size	Weight	Sweetness	Crunchiness	Juiciness	Ripeness	\
0	0.0	-3.970049	-2.512336	5.346330	-1.012009	1.844900	0.329840	
1	1.0	-1.195217	-2.839257	3.664059	1.588232	0.853286	0.867530	
2	2.0	-0.292024	-1.351282	-1.738429	-0.342616	2.838636	-0.038033	
3	3.0	-0.657196	-2.271627	1.324874	-0.097875	3.637970	-3.413761	
4	4.0	1.364217	-1.296612	-0.384658	-0.553006	3.030874	-1.303849	

	Acidity	Quality
0	-0.491590483	good
1	-0.722809367	good
2	2.621636473	bad
3	0.790723217	good
4	0.501984036	good

```
[65]: # converting categorical data to numerica
data['Quality']= data['Quality'].astype('category')
data['Quality']= data['Quality'].cat.codes
data.head()
```

```
[65]:
```

	A_id	Size	Weight	Sweetness	Crunchiness	Juiciness	Ripeness	\
0	0.0	-3.970049	-2.512336	5.346330	-1.012009	1.844900	0.329840	
1	1.0	-1.195217	-2.839257	3.664059	1.588232	0.853286	0.867530	
2	2.0	-0.292024	-1.351282	-1.738429	-0.342616	2.838636	-0.038033	
3	3.0	-0.657196	-2.271627	1.324874	-0.097875	3.637970	-3.413761	
4	4.0	1.364217	-1.296612	-0.384658	-0.553006	3.030874	-1.303849	

	Acidity	Quality
--	---------	---------

0	-0.491590483	1
1	-0.722809367	1
2	2.621636473	0
3	0.790723217	1
4	0.501984036	1

```
[66]: # findig where my data has null value
data.isnull().sum()
```

```
[66]: A_id      1
      Size      1
      Weight    1
      Sweetness 1
      Crunchiness 1
      Juiciness  1
      Ripeness   1
      Acidity     0
      Quality     0
      dtype: int64
```

```
[67]: #dropping the null data
data = data.dropna()
data.head()
```

```
[67]:   A_id      Size      Weight  Sweetness  Crunchiness  Juiciness  Ripeness  \
0    0.0 -3.970049 -2.512336   5.346330   -1.012009   1.844900   0.329840
1    1.0 -1.195217 -2.839257   3.664059    1.588232   0.853286   0.867530
2    2.0 -0.292024 -1.351282  -1.738429   -0.342616   2.838636  -0.038033
3    3.0 -0.657196 -2.271627   1.324874   -0.097875   3.637970  -3.413761
4    4.0  1.364217 -1.296612  -0.384658   -0.553006   3.030874  -1.303849

      Acidity  Quality
0  -0.491590483      1
1  -0.722809367      1
2   2.621636473      0
3   0.790723217      1
4   0.501984036      1
```

```
[68]: # splitting my data into independent and dependent variables
x= data.drop(columns='Quality')
x.head()
```

```
[68]:   A_id      Size      Weight  Sweetness  Crunchiness  Juiciness  Ripeness  \
0    0.0 -3.970049 -2.512336   5.346330   -1.012009   1.844900   0.329840
1    1.0 -1.195217 -2.839257   3.664059    1.588232   0.853286   0.867530
2    2.0 -0.292024 -1.351282  -1.738429   -0.342616   2.838636  -0.038033
3    3.0 -0.657196 -2.271627   1.324874   -0.097875   3.637970  -3.413761
```

```
4    4.0  1.364217 -1.296612 -0.384658    -0.553006    3.030874 -1.303849
```

```
    Acidity
0  -0.491590483
1  -0.722809367
2   2.621636473
3   0.790723217
4   0.501984036
```

```
[69]: #dependent variable
      y=data['Quality']
      y.head()
```

```
[69]: 0    1
      1    1
      2    0
      3    1
      4    1
      Name: Quality, dtype: int8
```

```
[70]: # Splitting the data into training and testing sets
      x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.2,
      ↪random_state=42)
```

```
[71]: # Ordinary Logistic Regression
      lr = LogisticRegression(max_iter=1000)
      lr.fit(x_train, y_train)
```

```
[71]: LogisticRegression(max_iter=1000)
```

```
[89]: # finding my y_pred and accuracy for ordinary#
      y_pred =lr.predict(x_test)
      accuracy_ordinary = accuracy_score(y_test,y_pred)
      print('y_pred:',y_pred)
      print('accuracy_ordinary',accuracy_ordinary)
```

```
y_pred: [1 1 1 0 0 1 0 0 0 0 1 1 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 1 1 0 1 0 0 1 1 0 1 0 0 0 1 1
0
1 0 0 0 0 0 1 0 1 1 1 1 0 1 0 1 1 1 0 0 0 1 1 0 1 0 1 0 1 0 0 1 0 1 1 1 0 1 0
1 1 1 0 0 1 0 1 0 1 0 0 0 0 1 1 0 0 1 1 0 0 1 0 1 1 0 1 0 0 0 0 1 0 0 0 1
0 0 0 0 0 0 0 0 1 0 0 1 1 0 0 0 1 1 1 0 1 1 1 1 1 0 0 1 0 1 1 1 0 1 0 0 0
0 1 0 0 0 1 1 0 1 0 0 0 0 1 0 0 1 0 0 1 0 1 0 1 0 1 0 0 1 0 1 0 1 0 0 0 0
1 1 0 0 0 1 1 0 0 0 1 0 1 0 1 0 0 1 1 0 0 0 0 0 1 0 1 1 0 1 1 1 1 1 1 1
1 0 0 1 1 0 0 1 1 1 1 1 1 1 1 0 1 1 0 1 1 1 0 0 0 0 0 0 1 0 0 1 0 1 0 0 0
0 0 0 0 0 0 0 1 1 0 0 0 1 0 0 0 1 0 1 0 0 1 0 0 0 1 0 0 0 0 0 1 0 1 0 0 1
0 1 0 0 0 1 1 1 1 1 1 1 1 0 1 1 1 1 0 0 0 1 1 1 0 1 1 1 0 1 1 1 0 0 0 0 1 0
0 1 1 0 0 0 0 0 0 1 0 0 1 1 0 1 1 0 1 0 0 1 1 1 1 0 1 1 1 0 1 1 1 0 0 1 1
1 1 1 1 0 0 1 1 0 0 1 1 1 1 1 0 0 1 1 0 1 0 1 0 1 0 1 1 0 1 0 1 0 1 1 1 1
```

```

1 0 1 0 1 0 1 1 0 0 1 1 1 1 1 0 1 1 1 1 1 0 1 0 1 0 1 0 0 0 1 1 1 1 0 1 1
1 0 1 0 1 0 0 1 0 0 0 0 1 1 0 0 0 1 1 0 1 0 1 1 1 0 0 0 1 1 1 0 0 1 0 1 0
0 0 1 0 1 0 1 1 0 0 0 0 0 0 1 0 1 1 0 1 1 1 1 0 1 1 1 1 0 1 0 1 1 0 0 1 1
1 1 1 1 1 1 1 0 0 1 1 0 1 0 1 1 1 0 1 1 0 0 1 1 0 1 0 0 1 1 0 0 0 1 1 0 1
1 0 0 0 0 1 0 0 1 1 1 1 1 0 1 0 1 1 0 1 0 1 0 1 0 0 0 1 0 1 1 0 0 0 1 1 1
1 1 1 0 1 1 1 1 0 1 1 0 0 0 1 0 0 0 0 0 0 1 1 0 0 0 0 0 0 0 1 0 0 1 0 0 1
0 1 1 0 1 0 0 1 0 1 0 1 1 1 0 0 0 1 1 1 0 0 1 1 1 1 0 1 0 0 1 0 1 1 1 1 0
1 0 0 1 0 0 1 1 0 0 0 1 0 0 0 1 1 1 0 0 1 0 0 1 0 1 1 1 1 0 1 0 0 0 0 1 0
0 1 1 0 1 0 0 0 0 1 1 1 0 1 0 1 0 1 1 0 1 1 0 0 0 1 1 0 0 0 0 1 0 0 1 1 0 0 1
1 1 0 1 0 1 1 0 0 1 0 1 0 0 0 0 0 0 0 1 1 0 1 1 1 1 1 1 0 0 0 1 1 0 1 1 0 0
1 0 0 1 1 1 1 1 0 0 1 0 0 1 0 0 0 0 1 1 1 1 1]
accuracy_ordinary 0.75625

```

```

[76]: #scaler
scaler = StandardScaler()
x_train_opt=scaler.fit_transform(x_train)
x_test_opt =scaler.transform(x_test)

```

```

[77]: # Building model for optimized LogisticRegression
lr_opt = LogisticRegression(random_state=0,
                             max_iter =1000,
                             C=1,
                             fit_intercept =True

                             ).fit(x_train,y_train).fit(x_train,y_train)

```

```

[90]: # finding my y_pred and accuracy for ordinar#
y_pred_opt =lr_opt.predict(x_test_opt)
accuracy_opt = accuracy_score(y_test,y_pred)
print('y_pred_opt:',y_pred_opt)
print('accuracy_opt',accuracy_opt)

```

```

y_pred_opt: [1 1 1 1 0 1 0 1 1 1 1 1 1 0 1 1 0 0 0 1 0 0 0 1 1 1 1 0 1 1 0 1 1 0 1
1 1 0
1 1 1 0 0 0 1 1 1 1 1 1 1 0 1 1 1 1 1 1 0 0 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1
1 1 1 0 0 1 0 1 0 1 1 1 0 1 1 1 1 0 1 1 1 1 1 0 1 1 1 1 1 0 0 0 1 0 1 1 1
1 0 0 1 0 1 0 1 1 0 0 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 0 1 1
1 1 0 0 1 1 1 0 1 1 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 0 1 1 1 0 0 1 1
1 1 0 0 0 1 1 0 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1
1 0 0 1 1 0 0 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 0 1 1 1 1 0 0 1 1 1 1 0 1 0 1 1
0 1 0 0 0 0 0 1 1 1 1 1 1 1 0 0 0 1 1 1 1 0 1 0 0 0 1 1 1 0 1 1 1 0 1 0 1 1
0 1 0 0 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 0 1 0 1 1
1 1 1 0 0 0 0 0 1 1 1 0 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1
1 1 1 1 1 0 1 1 0 0 1 1 1 1 1 0 0 1 1 1 1 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1
1 0 1 1 1 1 1 1 0 1 1 1 1 1 1 1 0 1 1 1 1 1 0 1 0 1 1 1 1 1 0 0 1 1 1 1 1 1 1
1 1 1 1 1 1 0 1 1 0 1 0 1 1 0 0 1 1 1 0 1 0 1 1 1 1 0 0 1 1 1 1 0 1 0 1 0
1 0 1 1 1 0 1 1 1 1 0 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 0 1 1
1 1 1 1 1 1 1 1 0 1 1 1 1 0 1 1 1 1 1 1 0 0 1 1 0 1 0 0 1 1 1 0 1 1 1 0 1

```

```

1 1 0 0 1 1 0 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 0 0 0 1 0 1 1 0 0 1 1 1 1
1 1 1 0 1 1 1 1 0 1 1 0 0 0 1 0 1 1 0 0 0 1 1 1 0 0 0 1 0 0 1 0 1 1 0 0 1
0 1 1 0 1 1 1 1 0 1 1 1 1 1 0 1 0 1 1 1 1 0 1 1 1 1 0 1 0 0 1 0 1 1 1 1 0
1 0 0 1 1 1 1 1 0 1 0 1 0 1 0 1 1 1 0 1 1 0 1 1 1 1 1 1 1 0 1 0 0 0 0 1 0
0 1 1 1 1 0 0 1 1 1 1 1 0 1 0 1 1 0 1 1 1 0 0 1 1 1 0 1 1 1 0 1 1 1 1 0 1
1 1 1 1 0 1 1 0 1 1 0 1 1 1 0 0 0 1 1 1 0 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 0 0
1 0 0 1 1 1 1 1 0 1 1 1 1 1 1 0 1 1 0 1 1 1 1 1]
accuracy_opt 0.75625

```

C:\Users\lynda\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names

```
warnings.warn(
```

```

[86]: if accuracy_ordinary>accuracy_opt:
        print('The ordinary model performs better than the optimized model')
elif accuracy_ordinary < accuracy_opt:
        print('The ordinary model performs better than the optimized')
else:
        print('Both modes have the same accuracy')

```

Both modes have the same accuracy

```
[ ]:
```

```
[ ]:
```