

Muhumuza Stephen 00526 linearregression

March 25, 2024

```
[51]: # importing the necessary libraries
import numpy as np
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import matplotlib.pyplot as plt
```

```
[52]: # Loading the dataset
data=pd.read_csv('C:\\Users\\lynda\\Desktop\\company.csv')
data.head()
```

```
[52]:
```

	R&D Spend	Administration	Marketing Spend	State	Profit
0	165349.20	136897.80	471784.10	New York	192261.83
1	162597.70	151377.59	443898.53	California	191792.06
2	153441.51	101145.55	407934.54	Florida	191050.39
3	144372.41	118671.85	383199.62	New York	182901.99
4	142107.34	91391.77	366168.42	Florida	166187.94

```
[53]: #checking for null values
data.isna().sum()
```

```
[53]:
```

R&D Spend	0
Administration	0
Marketing Spend	0
State	0
Profit	0

dtype: int64

```
[54]: # defining the x(independent valu)
x=data[['R&D Spend', 'Administration', 'Marketing Spend']]
x.head()
```

```
[54]:
```

	R&D Spend	Administration	Marketing Spend
0	165349.20	136897.80	471784.10
1	162597.70	151377.59	443898.53
2	153441.51	101145.55	407934.54

3	144372.41	118671.85	383199.62
4	142107.34	91391.77	366168.42

```
[55]: # y(dependent values)
y=data['Profit']
y.head()
```

```
[55]: 0    192261.83
      1    191792.06
      2    191050.39
      3    182901.99
      4    166187.94
      Name: Profit, dtype: float64
```

```
[56]: # splitting my data into training and testing set
x_train,x_test,y_train,y_test = train_test_split(x,y, test_size=0.2,
↳random_state=42)
```

```
[57]: # Building my Ordinary Linear regression
model = LinearRegression()
model.fit(x_train, y_train)
```

```
[57]: LinearRegression()
```

```
[58]: #calculating accuracy score
y_pred = model.predict(x_test)
y_pred
```

```
[58]: array([126703.02716461,  84894.75081556,  98893.41815974,  46501.70815036,
          129128.39734381,  50992.69486261, 109016.5536578 , 100878.4641454 ,
          97700.59638629, 113106.15292226])
```

```
[59]: r2=r2_score(y_test,y_pred)
r2
```

```
[59]: 0.900065308303732
```

```
[62]: # Optimization of Linear Regression
model_opt=LinearRegression()
model_opt
```

```
[62]: LinearRegression()
```

```
[87]: #define parameters
params={
    'fit_intercept': [True,False],
    'copy_X': [True,False],
    'n_jobs': [True,False],
```

```
    'positive': [True, False],  
}
```

```
[88]: #initializing the gridsearch  
grid_search=GridSearchCV(model_opt,params,cv=5)  
grid_search
```

```
[88]: GridSearchCV(cv=5, estimator=LinearRegression(),  
                  param_grid={'copy_X': [True, False],  
                              'fit_intercept': [True, False],  
                              'n_jobs': [True, False], 'positive': [True, False]})
```

```
[89]: # training optimized model_  
grid_search.fit(x_train,y_train)
```

```
[89]: GridSearchCV(cv=5, estimator=LinearRegression(),  
                  param_grid={'copy_X': [True, False],  
                              'fit_intercept': [True, False],  
                              'n_jobs': [True, False], 'positive': [True, False]})
```

```
[90]: #finding params  
best_params=grid_search.best_params_  
best_params
```

```
[90]: {'copy_X': True, 'fit_intercept': True, 'n_jobs': True, 'positive': True}
```

```
[91]: #finding the best model  
best_model=LinearRegression(**best_params)  
best_model
```

```
[91]: LinearRegression(n_jobs=True, positive=True)
```

```
[92]: #fitting model  
best_model.fit(x_train,y_train)
```

```
[92]: LinearRegression(n_jobs=True, positive=True)
```

```
[93]: #making predictions  
y_pred_opt=best_model.predict(x_test)  
y_pred_opt
```

```
[93]: array([127521.38604123,  82615.07411457,  97683.2462344 ,  46400.65677644,  
          130782.53611917,  45967.0205249 , 109813.19061887, 101612.68921418,  
          97023.64013854, 113241.36575804])
```

```
[94]: #evaluting my model  
r2=r2_score(y_test,y_pred_opt)
```

```
r2
```

```
[94]: 0.9168381183550245
```

```
[ ]:
```